

# LaonTap - Logistic Regression

LoanTap is an online platform committed to delivering customized loan products to millennials. They innovate in an otherwise dull loan segment, to deliver instant, flexible loans on consumer-friendly terms to salaried professionals and businessmen.

The data science team at LoanTap is building an underwriting layer to determine the creditworthiness of MSMEs as well as individuals.

**Problem Statement:** Given a set of attributes for an Individual, determine if a credit line should be extended to them. The main challenge is to minimise the risk of NPAs by flagging defaulters while maximising opportunity to earn interest by disbursing loans to as many customers as possible.

**Data dictionary:**

1. `loan_amnt` : The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.
2. `term` : The number of payments on the loan. Values are in months and can be either 36 or 60.
3. `int_rate` : Interest Rate on the loan
4. `installment` : The monthly payment owed by the borrower if the loan originates.

5. grade : LoanTap assigned loan grade
6. sub\_grade : LoanTap assigned loan subgrade
7. emp\_title : The job title supplied by the Borrower when applying for the loan.
8. emp\_length : Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years.
9. home\_ownership : The home ownership status provided by the borrower during registration or obtained from the credit report.
10. annual\_inc : The self-reported annual income provided by the borrower during registration.
11. verification\_status : Indicates if income was verified by LoanTap, not verified, or if the income source was verified
12. issue\_d : The month which the loan was funded
13. loan\_status : Current status of the loan - Target Variable
14. purpose : A category provided by the borrower for the loan request.
15. title : The loan title provided by the borrower
16. dti : A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, excluding mortgage and the requested LoanTap loan, divided by the borrower's self-reported monthly income.
17. earliest\_cr\_line : The month the borrower's earliest reported credit line was opened
18. open\_acc : The number of open credit lines in the borrower's credit file.
19. pub\_rec : Number of derogatory public records
20. revol\_bal : Total credit revolving balance
21. revol\_util : Revolving line utilization rate, or the amount of credit the borrower is using relative to all available revolving credit.
22. total\_acc : The total number of credit lines currently in the borrower's credit file
23. initial\_list\_status : The initial listing status of the loan. Possible values are - W, F
24. application\_type : Indicates whether the loan is an individual application or a joint application with two co-borrowers
25. mort\_acc : Number of mortgage accounts.
26. pub\_rec\_bankruptcies : Number of public record bankruptcies
27. Address: Address of the individual

## Importing Libraries

```
#Data processing
import pandas as pd
import numpy as np

#Data Visualisation
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
%matplotlib inline
```

```

#Setting option for full column view of Data
pd.set_option('display.max_columns', None)

#Stats & model building
from scipy import stats
from sklearn.preprocessing import OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import (accuracy_score, confusion_matrix,
                             roc_curve, auc, ConfusionMatrixDisplay,
                             f1_score, recall_score,
                             precision_score, precision_recall_curve,
                             average_precision_score,
                             classification_report)
from statsmodels.stats.outliers_influence import
variance_inflation_factor
from imblearn.over_sampling import SMOTE

#Hide warnings
import warnings
warnings.filterwarnings("ignore")

df = pd.read_csv("logistic_regression.csv")
df.head()

```

	loan_amnt	term	int_rate	installment	grade	sub_grade	\
0	10000.0	36 months	11.44	329.48	B	B4	
1	8000.0	36 months	11.99	265.68	B	B5	
2	15600.0	36 months	10.49	506.97	B	B3	
3	7200.0	36 months	6.49	220.65	A	A2	
4	24375.0	60 months	17.27	609.33	C	C5	

	emp_title	emp_length	home_ownership	annual_inc	\
0	Marketing	10+ years	RENT	117000.0	
1	Credit analyst	4 years	MORTGAGE	65000.0	
2	Statistician	< 1 year	RENT	43057.0	
3	Client Advocate	6 years	RENT	54000.0	
4	Destiny Management Inc.	9 years	MORTGAGE	55000.0	

	verification_status	issue_d	loan_status	purpose	\
0	Not Verified	Jan-2015	Fully Paid	vacation	
1	Not Verified	Jan-2015	Fully Paid	debt_consolidation	
2	Source Verified	Jan-2015	Fully Paid	credit_card	

3	Not Verified	Nov-2014	Fully Paid	credit_card
4	Verified	Apr-2013	Charged Off	credit_card

	title	dti	earliest	cr line	open acc	pub rec
\						
0	Vacation	26.24		Jun-1990	16.0	0.0
1	Debt consolidation	22.05		Jul-2004	17.0	0.0
2	Credit card refinancing	12.79		Aug-2007	13.0	0.0
3	Credit card refinancing	2.60		Sep-2006	6.0	0.0
4	Credit Card Refinance	33.95		Mar-1999	13.0	0.0

	revol_bal	revol_util	total_acc	initial_list_status
application_type \				
0	36369.0	41.8	25.0	w
INDIVIDUAL				
1	20131.0	53.3	27.0	f
INDIVIDUAL				
2	11987.0	92.2	26.0	f
INDIVIDUAL				
3	5472.0	21.5	13.0	f
INDIVIDUAL				
4	24584.0	69.8	43.0	f
INDIVIDUAL				

	mort_acc	pub_rec_bankruptcies
\		
0	0.0	0.0
1	3.0	0.0
2	0.0	0.0
3	0.0	0.0
4	1.0	0.0

	address
0	0174 Michelle Gateway\r\nMendozaberg, OK 22690
1	1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113
2	87025 Mark Dale Apt. 269\r\nNew Sabrina, WV 05113
3	823 Reid Ford\r\nDelacruzside, MA 00813
4	679 Luna Roads\r\nGreggshire, VA 11650

*#shape of data*

df.shape

(396030, 27)

```
# Statistical summary
```

```
df.describe()
```

	loan_amnt	int_rate	installment	annual_inc	\
count	396030.000000	396030.000000	396030.000000	3.960300e+05	
mean	14113.888089	13.639400	431.849698	7.420318e+04	
std	8357.441341	4.472157	250.727790	6.163762e+04	
min	500.000000	5.320000	16.080000	0.000000e+00	
25%	8000.000000	10.490000	250.330000	4.500000e+04	
50%	12000.000000	13.330000	375.430000	6.400000e+04	
75%	20000.000000	16.490000	567.300000	9.000000e+04	
max	40000.000000	30.990000	1533.810000	8.706582e+06	

	dti	open_acc	pub_rec	revol_bal	\
count	396030.000000	396030.000000	396030.000000	3.960300e+05	
mean	17.379514	11.311153	0.178191	1.584454e+04	
std	18.019092	5.137649	0.530671	2.059184e+04	
min	0.000000	0.000000	0.000000	0.000000e+00	
25%	11.280000	8.000000	0.000000	6.025000e+03	
50%	16.910000	10.000000	0.000000	1.118100e+04	
75%	22.980000	14.000000	0.000000	1.962000e+04	
max	9999.000000	90.000000	86.000000	1.743266e+06	

	revol_util	total_acc	mort_acc
pub_rec_bankruptcies			
count	395754.000000	396030.000000	358235.000000
	395495.000000		
mean	53.791749	25.414744	1.813991
	0.121648		
std	24.452193	11.886991	2.147930
	0.356174		
min	0.000000	2.000000	0.000000
	0.000000		
25%	35.800000	17.000000	0.000000
	0.000000		
50%	54.800000	24.000000	1.000000
	0.000000		
75%	72.900000	32.000000	3.000000
	0.000000		
max	892.300000	151.000000	34.000000
	8.000000		

## Data Cleaning

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
0   loan_amnt                            396030 non-null  float64
1   term                                396030 non-null  object
2   int_rate                            396030 non-null  float64
3   installment                         396030 non-null  float64
4   grade                               396030 non-null  object
5   sub_grade                          396030 non-null  object
6   emp_title                           373103 non-null  object
7   emp_length                         377729 non-null  object
8   home_ownership                     396030 non-null  object
9   annual_inc                         396030 non-null  float64
10  verification_status                396030 non-null  object
11  issue_d                           396030 non-null  object
12  loan_status                       396030 non-null  object
13  purpose                           396030 non-null  object
14  title                             394275 non-null  object
15  dti                               396030 non-null  float64
16  earliest_cr_line                  396030 non-null  object
17  open_acc                         396030 non-null  float64
18  pub_rec                         396030 non-null  float64
19  revol_bal                       396030 non-null  float64
20  revol_util                       395754 non-null  float64
21  total_acc                       396030 non-null  float64
22  initial_list_status              396030 non-null  object
23  application_type                 396030 non-null  object
24  mort_acc                       358235 non-null  float64
25  pub_rec_bankruptcies            395495 non-null  float64
26  address                          396030 non-null  object
dtypes: float64(12), object(15)
memory usage: 81.6+ MB

```

## Checking Column Datatypes

```

# Non-numeric columns
cat_cols = df.select_dtypes(include='object').columns
cat_cols

Index(['term', 'grade', 'sub_grade', 'emp_title', 'emp_length',
      'home_ownership', 'verification_status', 'issue_d',
      'loan_status',
      'purpose', 'title', 'earliest_cr_line', 'initial_list_status',
      'application_type', 'address'],
      dtype='object')

```

```

# Number of unique values in all non-numeric columns
for col in cat_cols:
    print(f"No. of unique values in {col}: {df[col].nunique()}")

No. of unique values in term: 2
No. of unique values in grade: 7
No. of unique values in sub_grade: 35
No. of unique values in emp_title: 173105
No. of unique values in emp_length: 11
No. of unique values in home_ownership: 6
No. of unique values in verification_status: 3
No. of unique values in issue_d: 115
No. of unique values in loan_status: 2
No. of unique values in purpose: 14
No. of unique values in title: 48817
No. of unique values in earliest_cr_line: 684
No. of unique values in initial_list_status: 2
No. of unique values in application_type: 3
No. of unique values in address: 393700

# Convert earliest credit line & issue date to datetime
df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'])
df['issue_d'] = pd.to_datetime(df['issue_d'])

# Convert employment length to numeric
d = {'10+ years':10, '4 years':4, '< 1 year':0,
      '6 years':6, '9 years':9, '2 years':2, '3 years':3,
      '8 years':8, '7 years':7, '5 years':5, '1 year':1}
df['emp_length'] = df['emp_length'].replace(d)

# Convert columns with less number of unique values to categorical
columns
cat_cols = ['term', 'grade', 'sub_grade', 'home_ownership',
            'verification_status', 'loan_status', 'purpose',
            'initial_list_status', 'application_type']

df[cat_cols] = df[cat_cols].astype('category')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 396030 entries, 0 to 396029
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
0   loan_amnt              396030 non-null  float64
1   term                   396030 non-null  category
2   int_rate                396030 non-null  float64
3   installment            396030 non-null  float64
4   grade                  396030 non-null  category
5   sub_grade              396030 non-null  category

```

```

6   emp_title      373103 non-null object
7   emp_length    377729 non-null float64
8   home_ownership 396030 non-null category
9   annual_inc     396030 non-null float64
10  verification_status 396030 non-null category
11  issue_d        396030 non-null datetime64[ns]
12  loan_status    396030 non-null category
13  purpose        396030 non-null category
14  title          394275 non-null object
15  dti            396030 non-null float64
16  earliest_cr_line 396030 non-null datetime64[ns]
17  open_acc       396030 non-null float64
18  pub_rec        396030 non-null float64
19  revol_bal      396030 non-null float64
20  revol_util     395754 non-null float64
21  total_acc      396030 non-null float64
22  initial_list_status 396030 non-null category
23  application_type 396030 non-null category
24  mort_acc       358235 non-null float64
25  pub_rec_bankruptcies 395495 non-null float64
26  address        396030 non-null object
dtypes: category(9), datetime64[ns](2), float64(13), object(3)
memory usage: 57.8+ MB

```

## Check for Duplicate Values

```
df.duplicated().sum()

0
```

There are no duplicate instances in the data

## Handling Missing Values

```
df.isna().sum()

loan_amnt      0
term           0
int_rate       0
installment    0
grade          0
sub_grade      0
emp_title      22927
emp_length     18301
home_ownership 0
annual_inc     0
verification_status 0
issue_d        0
loan_status    0

```



```
purpose          0
title            1755
dti              0
earliest_cr_line 0
open_acc         0
pub_rec          0
revol_bal        0
revol_util       276
total_acc        0
initial_list_status 0
application_type  0
mort_acc         37795
pub_rec_bankruptcies 535
address          0
dtype: int64
```

```
#Filling missing values with 'Unknown' for object dtype
fill_values = {'title': 'Unknown', 'emp_title': 'Unknown'}
df.fillna(value=fill_values, inplace=True)

#Mean aggregation of mort_acc by total_acc to fill missing values

avg_mort = df.groupby('total_acc')['mort_acc'].mean()

def fill_mort(total_acc, mort_acc):
    if np.isnan(mort_acc):
        return avg_mort[total_acc].round()
    else:
        return mort_acc

df['mort_acc'] = df.apply(lambda x:
    fill_mort(x['total_acc'],x['mort_acc']), axis=1)

df.dropna(inplace=True)

df.isna().sum()
```

```
loan_amnt        0
term             0
int_rate         0
installment      0
grade            0
sub_grade        0
emp_title        0
emp_length       0
home_ownership   0
annual_inc       0
verification_status 0
issue_d          0
loan_status      0
purpose          0
```

```
title          0
dti            0
earliest_cr_line 0
open_acc       0
pub_rec        0
revol_bal      0
revol_util     0
total_acc      0
initial_list_status 0
application_type 0
mort_acc       0
pub_rec_bankruptcies 0
address        0
dtype: int64

df.shape

(376929, 27)
```

## Outlier Treatment


```
num_cols = df.select_dtypes(include='number').columns
num_cols

Index(['loan_amnt', 'int_rate', 'installment', 'emp_length',
      'annual_inc',
      'dti', 'open_acc', 'pub_rec', 'revol_bal', 'revol_util',
      'total_acc',
      'mort_acc', 'pub_rec_bankruptcies'],
      dtype='object')


fig = plt.figure(figsize=(10,21))
i=1
for col in num_cols:
    ax = plt.subplot(7,2,i)
    sns.boxplot(x=df[col])
    plt.title(f'Boxplot of {col}')
    i += 1

plt.tight_layout()
plt.show()
```


Boxplot of loan\_amnt



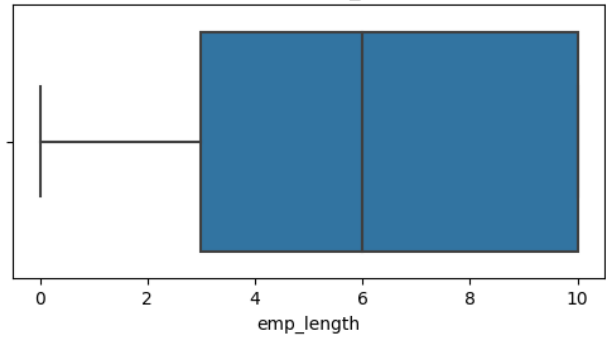
Boxplot of int\_rate




Boxplot of installment




Boxplot of emp\_length



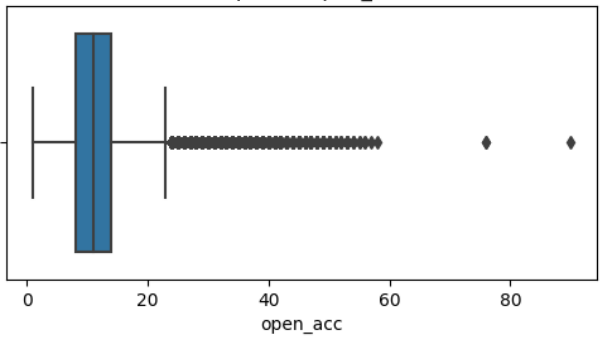
Boxplot of annual\_inc



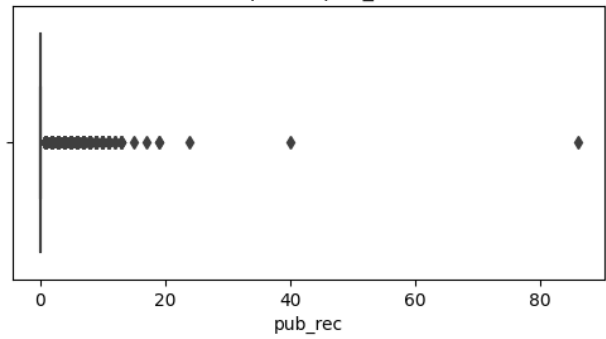
Boxplot of dti




Boxplot of open\_acc




Boxplot of pub\_rec



Boxplot of revol\_bal



Boxplot of revol\_util



Here we can see that many columns have outliers. Lets remove the rows with outliers using standard deviation (99% data is within 3 standard deviations in case of normally distributed data).

For pub\_Rec and pub\_rec\_bankruptcies, we can apply the 0 or 1 approach

```
# Convert pub_rec and pub_rec_bankruptcies to categorical variables

df['pub_rec_bankruptcies'] =
np.where(df['pub_rec_bankruptcies']>0, 'yes', 'no')
df['pub_rec'] = np.where(df['pub_rec']>0, 'yes', 'no')
df[['pub_rec_bankruptcies', 'pub_rec']] =
df[['pub_rec_bankruptcies', 'pub_rec']].astype('category')

# Numeric columns after converting public records to category
num_cols = df.select_dtypes(include='number').columns
num_cols

Index(['loan_amnt', 'int_rate', 'installment', 'emp_length',
      'annual_inc',
      'dti', 'open_acc', 'revol_bal', 'revol_util', 'total_acc',
      'mort_acc'],
      dtype='object')

#Removing outliers using standard deviation
for col in num_cols:
    mean=df[col].mean()
    std=df[col].std()
    upper = mean + (3*std)
    df = df[~(df[col]>upper)]

df.shape

(350845, 27)
```

## Feature Engineering

```
df['address'].sample(10)

285569          Unit 9894 Box 9319\r\nDPO AA 05113
101576    2867 Lindsey Shoal\r\nWilliamschester, LA 00813
139688    008 Alicia Gateway\r\nLake Stacey, VT 30723
160700    2315 Pamela Park\r\nNew Aaronbury, HI 05113
219251    3828 Jack Squares Suite 231\r\nRodriguezhaven,...
373303    0136 Tina Inlet\r\nNew Frankton, MO 30723
299038    949 Adam Track\r\nNorth Ryanberg, HI 00813
109538    12924 White Island\r\nLisaborough, WI 30723
360921    0203 Keith Neck\r\nEast Brandon, SC 30723
319103    607 Jennifer Path\r\nNew Williamtown, HI 48052
Name: address, dtype: object
```

```
# Deriving zip code and state from address
df[['state', 'zip_code']] = df['address'].apply(lambda x:
pd.Series([x[-8:-6], x[-5:]])

#Drop address
df.drop(["address"], axis = 1, inplace=True)

df.zip_code.nunique()


10
```

Since there are only 10 zipcodes, we can change the datatype of zipcodes to categorical

```
df['zip_code'] = df['zip_code'].astype('category')
```

## Exploratory Data Analysis

```
#Correlation between numerical features
plt.figure(figsize=(8,6))
sns.heatmap(df.corr(), annot=True, fmt=".1f")
plt.title('Correlation between Numerical Features')
plt.show()
```



1. loan\_amnt and installment are perfectly correlated
2. total\_acc is highly correlated with open\_acc
3. total\_acc is moderately correlated with mort\_acc We can remove some of these correlated features to avoid multicollinearity



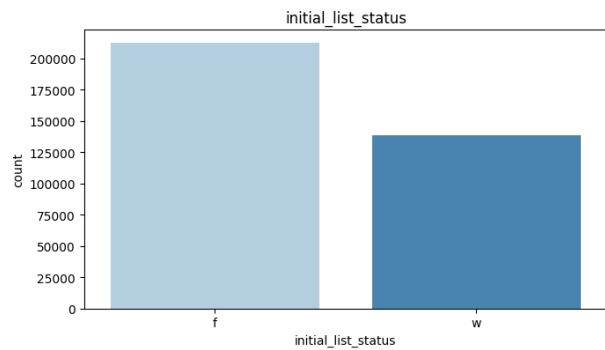
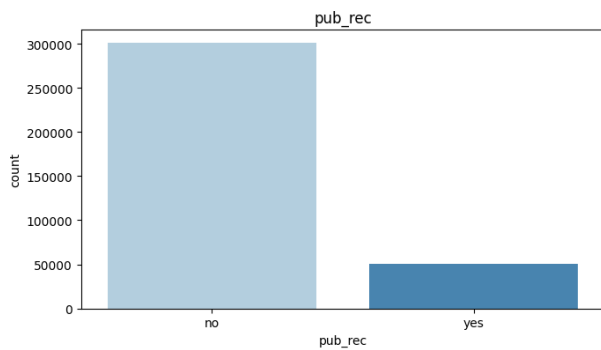

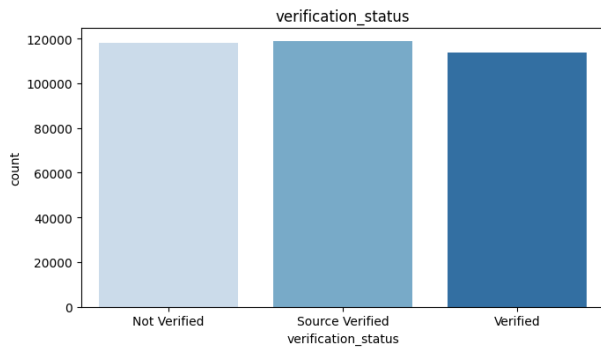
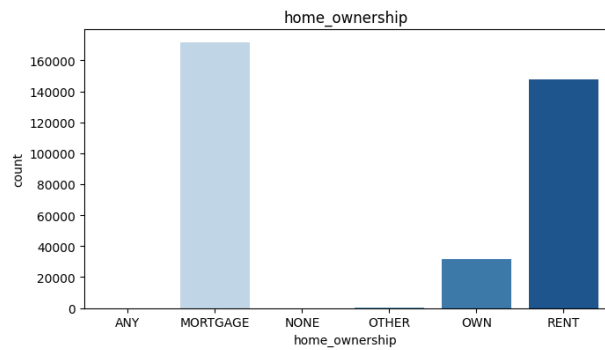
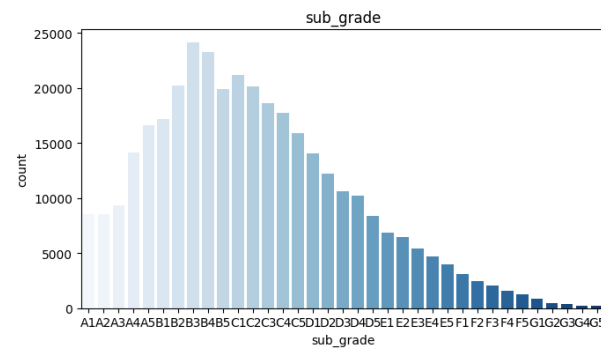


```
#Drop installment
df.drop(columns=['installment'], inplace=True)

#Distribution of categorical variables
plot = ['term', 'grade', 'sub_grade', 'home_ownership',
        'verification_status',
        'loan_status', 'pub_rec', 'initial_list_status',
        'application_type', 'pub_rec_bankruptcies']

plt.figure(figsize=(14,20))
i=1
for col in plot:
    ax=plt.subplot(5,2,i)
```

```
sns.countplot(x=df[col], palette='Blues')
plt.title(f'{col}')
i += 1

plt.tight_layout()
plt.show()
```

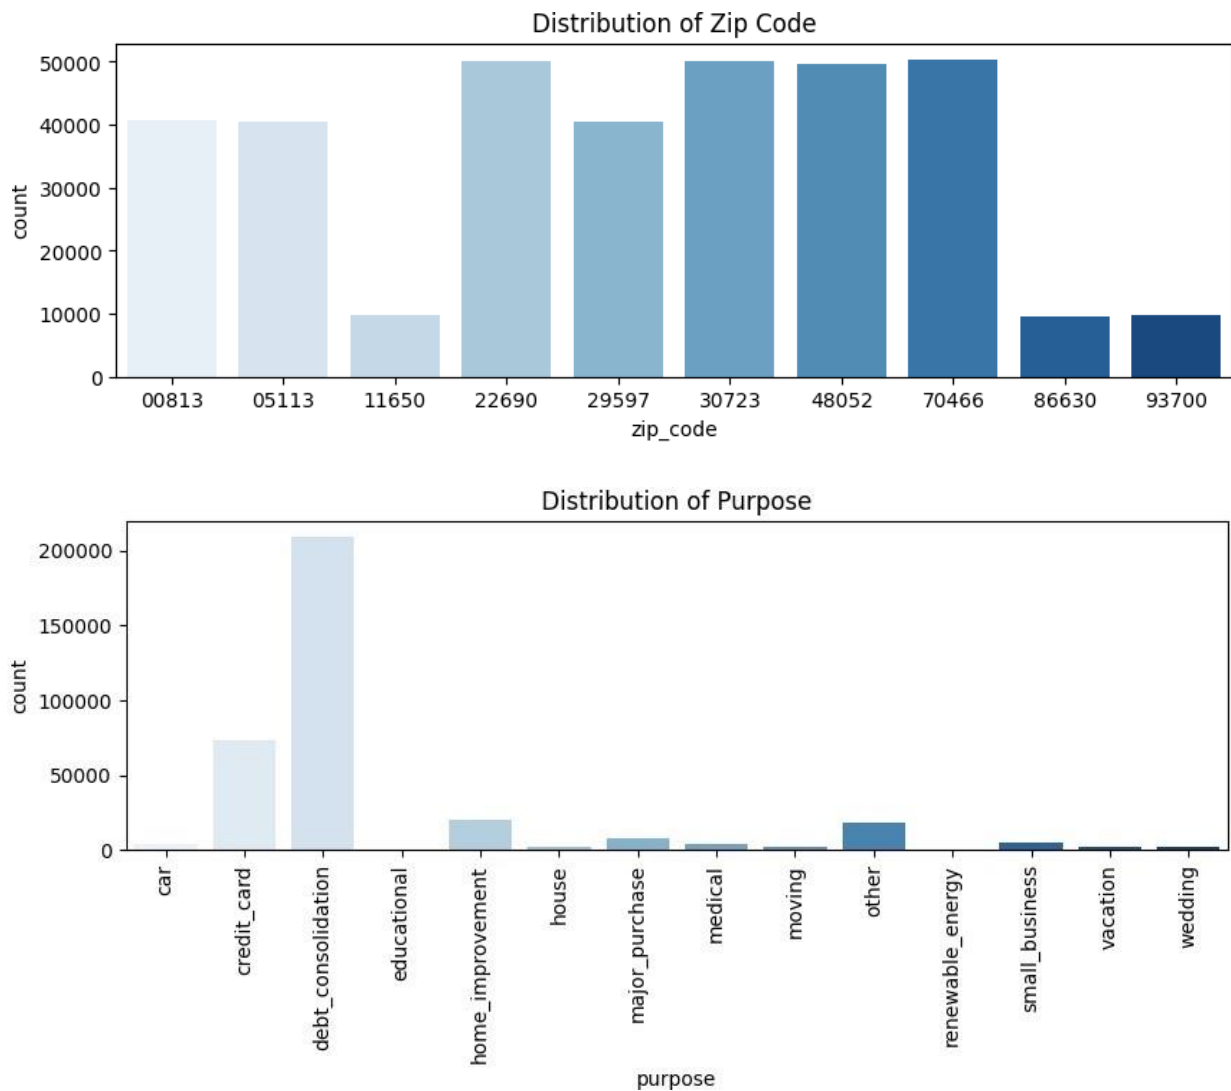




```
plt.figure(figsize=(10,3))
sns.countplot(x=df['zip_code'], palette='Blues')
plt.title('Distribution of Zip Code')

plt.figure(figsize=(10,3))
sns.countplot(x=df['purpose'], palette='Blues')
plt.xticks(rotation=90)
plt.title('Distribution of Purpose')

plt.show()
```



Observations:

- Almost 80% loans are of 36 months term
- Maximum loans (30%) fall in B grade, followed by C,A & D respectively
- The type of home ownership for 50% cases is mortgage

- The target variable (loan status) is imbalanced in the favour of fully-paid loans. Defaulters are approx 25% of fully paid instances.
- 85% of applicants don't have a public record/haven't filled for bankruptcy
- 99% applicants have applied under 'individual' application type
- 55% of loans are taken for the purpose of debt consolidation followed by 20% on credit card

```
# Impact of categorical factors on loan status



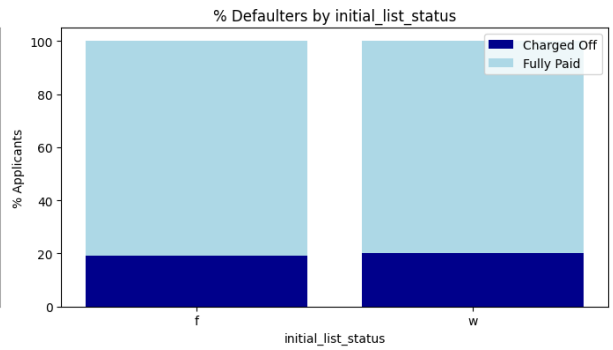
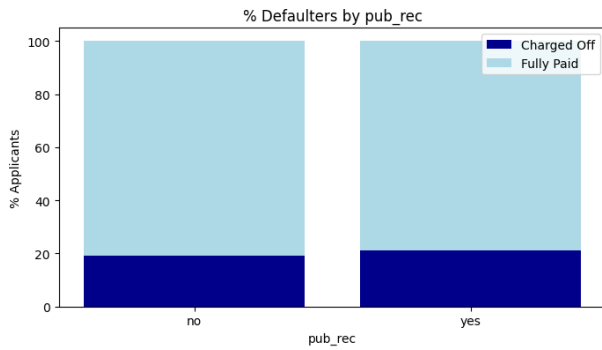


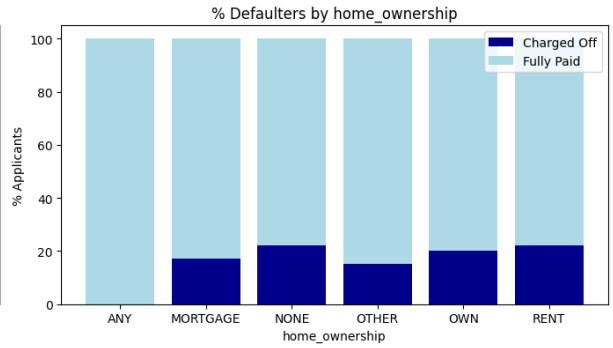


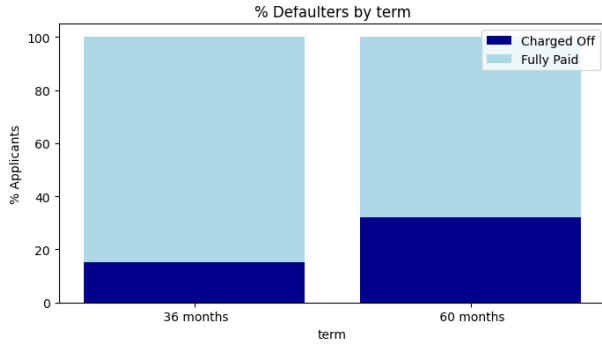
plot = ['term', 'grade', 'sub_grade', 'home_ownership',
        'verification_status',
        'zip_code', 'pub_rec', 'initial_list_status',
        'application_type', 'pub_rec_bankruptcies']

plt.figure(figsize=(14,20))
i=1
for col in plot:
    ax=plt.subplot(5,2,i)

    data = df.pivot_table(index=col, columns='loan_status',
aggfunc='count', values='purpose')
    data = data.div(data.sum(axis=1), axis=0).multiply(100).round()
    data.reset_index(inplace=True)

    plt.bar(data[col],data['Charged Off'], color='#00008b')
    plt.bar(data[col],data['Fully Paid'], color='#add8e6',
bottom=data['Charged Off'])
    plt.xlabel(f'{col}')
    plt.ylabel('% Applicants')
    plt.title(f'% Defaulters by {col}')
    plt.legend(['Charged Off','Fully Paid'])
    i += 1

plt.tight_layout()
plt.show()
```



```

# Impact of Purpose/state on loan status


purpose = df.pivot_table(index='purpose', columns='loan_status',
aggfunc='count', values='sub_grade')
purpose = purpose.div(purpose.sum(axis=1),
axis=0).multiply(100).round()
purpose.reset_index(inplace=True)


plt.figure(figsize=(14,4))
plt.bar(purpose['purpose'],purpose['Charged Off'], color='#00008b')
plt.bar(purpose['purpose'],purpose['Fully Paid'], color='#add8e6',
bottom=purpose['Charged Off'])
plt.xlabel('Purpose')
plt.ylabel('% Applicants')
plt.title('% Defaulters by purpose')
plt.legend(['Charged Off','Fully Paid'])
plt.xticks(rotation=90)
plt.show()

state = df.pivot_table(index='state', columns='loan_status',
aggfunc='count', values='sub_grade')
state = state.div(state.sum(axis=1), axis=0).multiply(100).round()
state.reset_index(inplace=True)

plt.figure(figsize=(14,4))
plt.bar(state['state'],state['Charged Off'], color='#00008b')
plt.bar(state['state'],state['Fully Paid'], color='#add8e6',
bottom=state['Charged Off'])
plt.xlabel('state')
plt.ylabel('% Applicants')
plt.title('% Defaulters by state')
plt.legend(['Charged Off','Fully Paid'])
plt.show()

```





### Observations:











- The % of defaulters is much higher for longer (60-month) term
- As expected, grade/sub-grade has the maximum impact on loan\_status with highest grade having maximum defaulters
- Zip codes such as 11650, 86630 and 93700 have 100% defaulters
- We can remove initial\_list\_status and state as they have no impact on loan\_status
- public records also don't seem to have any impact on loan\_status surprisingly
- Direct pay application type has higher default rate compared to individual/joint
- Loan taken for the purpose of small business has the highest rate of default

```
# Impact of numerical features on loan_status

num_cols = df.select_dtypes(include='number').columns

fig, ax = plt.subplots(10, 2, figsize=(15, 40))
i=0
color_dict = {'Fully Paid': matplotlib.colors.to_rgba('#add8e6', 0.5),
              'Charged Off': matplotlib.colors.to_rgba('#00008b', 1)}
for col in num_cols:
    sns.histplot(data=df, x=col, hue='loan_status', ax=ax[i, 0],
                 legend=True,
                 palette=color_dict, kde=True, fill=True)
    sns.boxplot(data=df, y=col, x='loan_status', ax=ax[i, 1],
                palette=('#00008b', '#add8e6'))
    ax[i, 0].set_ylabel(col, fontsize=12)
    ax[i, 0].set_xlabel(' ')
    ax[i, 1].set_xlabel(' ')
    ax[i, 1].set_ylabel(' ')
    ax[i, 1].xaxis.set_tick_params(labelsize=14)
    i += 1

plt.tight_layout()
plt.show()
```



## Observations:

- From the boxplots, it can be observed that the mean loan\_amnt, int\_rate, dti, open\_acc and revol\_util are slightly higher for defaulters while annual income is lower

```
# Remove columns which do not have an impact on loan_status
df.drop(columns=['initial_list_status','state',
                 'emp_title', 'title','earliest_cr_line',
                 'issue_d','sub_grade'], inplace=True)

# Subgrade is removed because grade and subgrade are similar features
```

## Data Pre-Processing

```
# Encoding Target Variable

df['loan_status']=df['loan_status'].map({'Fully Paid': 0, 'Charged
Off':1}).astype(int)

x = df.drop(columns=['loan_status'])
x.reset_index(inplace=True, drop=True)
y = df['loan_status']
y.reset_index(drop=True, inplace=True)

# Encoding Binary features into numerical dtype

x['term']=x['term'].map({' 36 months': 36, ' 60
months':60}).astype(int)
x['pub_rec']=x['pub_rec'].map({'no': 0, 'yes':1}).astype(int)
x['pub_rec_bankruptcies']=x['pub_rec_bankruptcies'].map({'no': 0,
'yes':1}).astype(int)
```

## One Hot Encoding of Categorical Features

```
cat_cols = x.select_dtypes('category').columns

encoder = OneHotEncoder(sparse=False)
encoded_data = encoder.fit_transform(x[cat_cols])
encoded_df = pd.DataFrame(encoded_data,
columns=encoder.get_feature_names_out(cat_cols))
x = pd.concat([x,encoded_df], axis=1)
x.drop(columns=cat_cols, inplace=True)
x.head()
```

	loan amnt	term	int rate	emp length	annual inc	dti	open acc
0	10000.0	36	11.44	10.0	117000.0	26.24	16.0
1	8000.0	36	11.99	4.0	65000.0	22.05	17.0

2	15600.0	36	10.49	0.0	43057.0	12.79	13.0
3	7200.0	36	6.49	6.0	54000.0	2.60	6.0
4	24375.0	60	17.27	9.0	55000.0	33.95	13.0

	pub_rec	revo	l_bal	revol_util	total_acc	mort_acc
pub_rec_bankruptcies \						
0	0	36	369.0	41.8	25.0	0.0
0						
1	0	20	31.0	53.3	27.0	3.0
0						
2	0	11	87.0	92.2	26.0	0.0
0						
3	0	5	72.0	21.5	13.0	0.0
0						
4	0	24	84.0	69.8	43.0	1.0
0						

	grade_A	grade_B	grade_C	grade_D	grade_E	grade_F	grade_G	\
0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
2	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
3	1.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	0.0	1.0	0.0	0.0	0.0	0.0	

	home_ownership_ANY	home_ownership_MORTGAGE	home_ownership_NONE	\
0	0.0	0.0	0.0	
1	0.0	1.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	1.0	0.0	

	home_ownership_OTHER	home_ownership_OWN	home_ownership_RENT	\
0	0.0	0.0	1.0	
1	0.0	0.0	0.0	
2	0.0	0.0	1.0	
3	0.0	0.0	1.0	
4	0.0	0.0	0.0	

	verification_status_Not Verified	verification_status_Source
Verified \		
0	1.0	
0.0		
1	1.0	
0.0		
2	0.0	
1.0		



3	1.0
0.0	
4	0.0
0.0	

	verification_status_Verified	purpose_car	purpose_credit_card \
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	1.0
3	0.0	0.0	1.0
4	1.0	0.0	1.0

	purpose_debt_consolidation	purpose_ed ucational
purpose_home_improvement \		
0	0.0	0.0
0.0		
1	1.0	0.0
0.0		
2	0.0	0.0
0.0		
3	0.0	0.0
0.0		
4	0.0	0.0
0.0		

	purpose_house	purpose_major_purchase	purpose_medical
purpose_moving \			
0	0.0	0.0	0.0
0.0			
1	0.0	0.0	0.0
0.0			
2	0.0	0.0	0.0
0.0			
3	0.0	0.0	0.0
0.0			
4	0.0	0.0	0.0
0.0			

	purpose_other	purpose_renewable_energy	purpose_small_business
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

	purpose_vacation	purpose_wedding	application_type_DIRECT_PAY \
0	1.0	0.0	0.0
1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0

4	0.0	0.0	0.0
	application_type_INDIVIDUAL	application_type_JOINT	zip_code_00813
\			
0	1.0	0.0	0.0
1	1.0	0.0	0.0
2	1.0	0.0	0.0
3	1.0	0.0	1.0
4	1.0	0.0	0.0

	zip_code_05113	zip_code_11650	zip_code_22690	zip_code_29597	\
0	0.0	0.0	1.0	0.0	
1	1.0	0.0	0.0	0.0	
2	1.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	1.0	0.0	0.0	

	zip_code_30723	zip_code_48052	zip_code_70466	zip_code_86630	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	zip_code_93700
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

## Train-Test Split

```
x_train, x_test, y_train, y_test =
train_test_split(x,y,test_size=0.20,stratify=y,random_state=42)

x_train.shape, y_train.shape, x_test.shape, y_test.shape

((280676, 56), (280676,), (70169, 56), (70169,))
```

## Scaling Numeric Features

```
scaler = MinMaxScaler()
x_train = pd.DataFrame(scaler.fit_transform(x_train),
columns=x_train.columns)
```

```
x_test = pd.DataFrame(scaler.transform(x_test),
columns=x_test.columns)
```

```
x_train.tail()
```

	loan_amnt	term	int_rate	emp_length	annual_inc	dti
open_acc \						
280671	0.167959	0.0	0.141671	0.7	0.194444	0.255954
0.60						
280672	0.497416	0.0	0.445778	0.4	0.182540	0.414482
0.24						
280673	0.064599	0.0	0.686664	0.7	0.238095	0.220111
0.32						
280674	0.245478	1.0	0.177665	0.9	0.313492	0.134953
0.92						
280675	0.646641	1.0	0.885095	0.6	0.349206	0.747173
0.88						

	pub_rec	revol_bal	revol_util	total_acc	mort_acc
280671	0.0	0.104275	0.271695	0.578947	0.428571
280672	0.0	0.224536	0.670722	0.263158	0.285714
280673	0.0	0.249454	0.622871	0.385965	0.428571
280674	0.0	0.080701	0.039740	0.842105	0.428571
280675	1.0	0.213775	0.543390	0.596491	0.714286

	pub_rec_bankruptcies	grade_A	grade_B	grade_C	grade_D
grade_E \					
280671	0.0	1.0	0.0	0.0	0.0
0.0					
280672	0.0	0.0	0.0	1.0	0.0
0.0					
280673	0.0	0.0	0.0	0.0	1.0
0.0					
280674	0.0	0.0	1.0	0.0	0.0
0.0					
280675	1.0	0.0	0.0	0.0	0.0
0.0					

	grade_F	grade_G	home_ownership_ANY	home_ownership_MORTGAGE
\				
280671	0.0	0.0	0.0	0.0
280672	0.0	0.0	0.0	0.0
280673	0.0	0.0	0.0	1.0
280674	0.0	0.0	0.0	1.0
280675	1.0	0.0	0.0	1.0

	home_ownership_NONE	home_ownership_OTHER	home_ownership_OWN
\			
280671	0.0	0.0	0.0
280672	0.0	0.0	0.0
280673	0.0	0.0	0.0
280674	0.0	0.0	0.0
280675	0.0	0.0	0.0
	home_ownership_RENT	verification_status_Not Verified	\
280671	1.0	1.0	
280672	1.0	0.0	
280673	0.0	0.0	
280674	0.0	0.0	
280675	0.0	0.0	
	verification_status_Source Verified		
verification_status_Verified \			
280671	0.0		
0.0			
280672	0.0		
1.0			
280673	1.0		
0.0			
280674	0.0		
1.0			
280675	0.0		
1.0			
	purpose_car	purpose_credit_card	
purpose_debt_consolidation \			
280671	0.0	1.0	0.0
280672	0.0	0.0	1.0
280673	0.0	0.0	0.0
280674	0.0	0.0	0.0
280675	0.0	0.0	0.0
	purpose_educational	purpose_home_improvement	
purpose house \			
280671	0.0	0.0	0.0
280672	0.0	0.0	0.0

280673	0.0	0.0	0.0
280674	0.0	0.0	0.0
280675	0.0	1.0	0.0

	purpose_major_purchase	purpose_medical	purpose_moving	\
280671	0.0	0.0	0.0	
280672	0.0	0.0	0.0	
280673	0.0	0.0	1.0	
280674	0.0	0.0	0.0	
280675	0.0	0.0	0.0	

	purpose_other	purpose_renewable_energy
purpose_small_business	\	
280671	0.0	0.0
0.0		
280672	0.0	0.0
0.0		
280673	0.0	0.0
0.0		
280674	0.0	1.0
0.0		
280675	0.0	0.0
0.0		

	purpose_vacation	purpose_wedding	application_type_DIRECT_PAY
\			
280671	0.0	0.0	0.0
280672	0.0	0.0	0.0
280673	0.0	0.0	0.0
280674	0.0	0.0	0.0
280675	0.0	0.0	0.0

	application_type_INDIVIDUAL	application_type_JOINT
zip_code_00813	\	
280671	1.0	0.0
0.0		
280672	1.0	0.0
1.0		
280673	1.0	0.0
0.0		
280674	1.0	0.0

1.0				
280675		1.0		0.0
0.0				
	zip_code_05113	zip_code_11650	zip_code_22690	zip_code_29597
\				
280671	0.0	0.0	0.0	0.0
280672	0.0	0.0	0.0	0.0
280673	0.0	0.0	0.0	0.0
280674	0.0	0.0	0.0	0.0
280675	0.0	1.0	0.0	0.0
	zip_code_30723	zip_code_48052	zip_code_70466	zip_code_86630
\				
280671	0.0	1.0	0.0	0.0
280672	0.0	0.0	0.0	0.0
280673	0.0	1.0	0.0	0.0
280674	0.0	0.0	0.0	0.0
280675	0.0	0.0	0.0	0.0
	zip_code_93700			
280671	0.0			
280672	0.0			
280673	0.0			
280674	0.0			
280675	0.0			

## Oversampling with SMOTE

```
# Oversampling to balance the target variable

sm=SMOTE(random_state=42)
x_train_res, y_train_res = sm.fit_resample(x_train,y_train.ravel())

print(f"Before OverSampling, count of label 1: {sum(y_train == 1)}")
print(f"Before OverSampling, count of label 0: {sum(y_train == 0)}")
print(f"After OverSampling, count of label 1: {sum(y_train_res == 1)}")
print(f"After OverSampling, count of label 0: {sum(y_train_res == 0)}")
```

```
Before OverSampling, count of label 1: 54200
Before OverSampling, count of label 0: 226476
After OverSampling, count of label 1: 226476
After OverSampling, count of label 0: 226476
```

## Logistic Regression

```
model = LogisticRegression()
model.fit(x_train_res, y_train_res)
train_preds = model.predict(x_train)
test_preds = model.predict(x_test)


#Model Evaluation
print('Train Accuracy:', model.score(x_train, y_train).round(2))
print('Train F1 Score:', f1_score(y_train, train_preds).round(2))
print('Train Recall
Score:', recall_score(y_train, train_preds).round(2))
print('Train Precision
Score:', precision_score(y_train, train_preds).round(2))

print('\nTest Accuracy:', model.score(x_test, y_test).round(2))
print('Test F1 Score:', f1_score(y_test, test_preds).round(2))
print('Test Recall Score:', recall_score(y_test, test_preds).round(2))
print('Test Precision
Score:', precision_score(y_test, test_preds).round(2))

# Confusion Matrix
cm = confusion_matrix(y_test, test_preds)
disp = ConfusionMatrixDisplay(cm)
disp.plot()
plt.title('Confusion Matrix')
plt.show()
```

```
Train Accuracy : 0.8
Train F1 Score: 0.61
Train Recall Score: 0.81
Train Precision Score: 0.49
```

```
Test Accuracy : 0.8
Test F1 Score: 0.61
Test Recall Score: 0.81
Test Precision Score: 0.49
```



## Model Coefficient with column name

```
coef = model.coef_[0]
imp_features = pd.DataFrame({"Features" : cols , "Weights": coef})
imp_features = imp_features.sort_values(["Weights"] , ascending = False)
print(model.coef_)
print(imp_features)
```

```
[[ 0.67086675  0.46294261  0.52257871 -0.01778102 -1.63378648  1.06860903
  0.84908845  0.05776276 -0.58046684  0.60909815 -0.7658847 -0.3039249
 -0.05502705 -0.97401973 -0.50760769 -0.11422166  0.09561079  0.23385738
  0.27905452  0.32370805 -0.02843792 -0.20084565  0.06421717 -0.31565695
 -0.15167948 -0.03121551 -0.29875822 -0.09959339 -0.26526672 -0.43969934
  0.02089984  0.10007305  0.21091204  0.16489038 -0.19889648 -0.02792455
 -0.0150925 -0.0195308  0.11368533 -0.0129795  0.50502374 -0.2612935
 -0.80368605 -0.35193795  0.66377059 -0.97545098 -8.53933083 -8.52879241
  8.71237456 -0.28498983 -8.52156752 -0.29044261 -0.2508412 -0.28609553
  8.58983296  8.73623407]]
```



```

-----
              Features  Weights
55          zip_code_93700  8.736234
48          zip_code_11650  8.712375
54          zip_code_86630  8.589833
5              dti  1.068609
6          open_acc  0.849088
0          loan_amnt  0.670867
44  application_type_INDIVIDUAL  0.663771
9              revol_util  0.609098
2              int_rate  0.522579
40  purpose_small_business  0.505024
1              term  0.462943
19              grade_G  0.323708
18              grade_F  0.279055
17              grade_E  0.233857
32  purpose_educational  0.210912
33  purpose_home_improvement  0.164890
38  purpose_other  0.113685
31  purpose_debt_consolidation  0.100073
16              grade_D  0.095611
22  home_ownership_NONE  0.064217
7              pub_rec  0.057763
30  purpose_credit_card  0.020900
39  purpose_renewable_energy -0.012979
36  purpose_medical -0.015093
3              emp_length -0.017781
37  purpose_moving -0.019531
35  purpose_major_purchase -0.027925
20  home_ownership_ANY -0.028438
25  home_ownership_RENT -0.031216
12  pub_rec_bankruptcies -0.055027
27  verification_status_Source Verified -0.099593
15              grade_C -0.114222
24  home_ownership_OWEN -0.151679

```

## Classification Report

```
print(classification_report(y_test, test_preds))
```

	precision	recall	f1-score	support
0	0.95	0.80	0.87	56619

1	0.49	0.81	0.61	13550
accuracy			0.80	70169
macro avg	0.72	0.80	0.74	70169
weighted avg	0.86	0.80	0.82	70169

- It can be observed that the recall score is very high (our model is able to identify 80% of actual defaulters) but the precision is low for positive class (of all the predicted defaulters, only 50% are actually defaulters).
- Although this model is effective in reducing NPAs by flagging most of the defaulters, it may cause loantap to deny loans to many deserving customers due to low precision (false positives)
- Low precision has also caused F1 score to drop to 60% even though accuracy is 80%


## Feature Importance

```

feature_imp = pd.DataFrame({'Columns':x_train.columns,
'Coefficients':model.coef_[0]}).round(2).sort_values('Coefficients',
ascending=False)

plt.figure(figsize=(8,8))
sns.barplot(y = feature_imp['Columns'],
            x = feature_imp['Coefficients'])
plt.title("Feature Importance for Model")
plt.yticks(fontsize=8)
plt.ylabel("Feature")
plt.tight_layout()
plt.show()

```



- The model has assigned large weightage to zip\_code features followed by dti, open\_acc, loan\_amnt
- Similarly, large negative coefficients are assigned to a few zip codes, followed by annual income and joint application type

## ROC Curve & AUC

The Receiver Operating Characteristic (ROC) curve is a graphical representation of the performance of a binary classification model. It helps evaluate and compare different models by illustrating the trade-off between the true positive rate (TPR) and false positive rate (FPR) at various classification thresholds.

The ROC curve is created by plotting the TPR on the y-axis against the FPR on the x-axis for different threshold values.

- TPR: Also known as sensitivity or recall, is the proportion of true positive predictions out of all actual positive instances.
- FPR: Proportion of false positive predictions out of all actual negative instances.

A perfect classifier would have a TPR of 1 and an FPR of 0, resulting in a point at the top-left corner of the ROC curve. On the other hand, a random classifier would have an ROC curve following the diagonal line, as it has an equal chance of producing true positive and false positive predictions.

The area under the ROC curve (AUC) is a commonly used metric to quantify the overall performance of a classifier.

A perfect classifier would have an AUC of 1, while a random classifier would have an AUC of 0.5. The higher the AUC value, the better the classifier's performance in distinguishing between positive and negative instances.


```
# Predict probabilities for the test set
probs = model.predict_proba(x_test)[: ,1]

# Compute the false positive rate, true positive rate, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, probs)

# Compute the area under the ROC curve
roc_auc = auc(fpr, tpr)

# Plot the ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
```

```
plt.legend(loc="lower right")
plt.show()
```



- AUC of 0.91 signifies that the model is able to discriminate well between the positive and the negative class.
- But it is not a good measure for an imbalanced target variable because it may be high even when the classifier has a poor score on the minority class.
- This can happen when the classifier performs well on the majority class instances, which dominate the dataset. As a result, the AUC may appear high, but the model may not effectively identify the minority class instances.

Lets plot the Precision-Recall curve which is more suited for evaluation of imbalanced data

### Precision Recall Curve

The Precision-Recall (PR) curve is another graphical representation commonly used to evaluate the performance of a binary classification model. It provides insights into the trade-off between precision and recall at various classification thresholds.

- Precision represents the proportion of correctly predicted positive instances out of all instances predicted as positive. It focuses on the accuracy of positive predictions.

- Recall, also known as sensitivity or true positive rate, represents the proportion of correctly predicted positive instances out of all actual positive instances. It focuses on capturing all positive instances.

Similar to the ROC curve, the PR curve is created by plotting recall on the x-axis and precision on the y-axis for different threshold values. The curve illustrates the relationship between precision and recall as the classification threshold changes.

A perfect classifier would have a precision of 1 and a recall of 1, resulting in a point at the top-right corner of the PR curve. Conversely, a random classifier would have a PR curve following the horizontal line defined by the ratio of positive instances in the dataset.


The PR curve is useful when dealing with imbalanced datasets, where the number of negative instances far outweighs the positives. In such cases, the PR curve provides a more comprehensive evaluation of the model's performance compared to the ROC curve. This is because the ROC curve can be misleading when the majority of instances are negative, as it primarily focuses on the true negative rate.

The area under the PR curve (AUPRC) is a commonly used metric to quantify the overall performance of a classifier. A perfect classifier would have an AUPRC of 1, while a random classifier would have an AUPRC equal to the ratio of positive instances. Generally, a higher AUPRC indicates better performance.

```
# Compute the false precision and recall at all thresholds
precision, recall, thresholds = precision_recall_curve(y_test, probs)

# Area under Precision Recall Curve
auprc = average_precision_score(y_test, probs)

# Plot the precision-recall curve
plt.plot(recall, precision, marker='.', label='PR curve (area = %0.2f)' % auprc)
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.title('Precision-Recall Curve')
plt.legend(loc="lower left")
plt.show()
```



As expected, the area under precision recall curve is not as high. It is a decent model as the area is more than 0.5 (random model benchmark) but there is still scope for improvement

## Tradeoff Questions

1. How can we make sure that our model can detect real defaulters and there are less false positives? This is important as we can lose out on an opportunity to finance more individuals and earn interest on it.
  - Answer - Since data is imbalances by making the data balance, we can try to avoid false positives. For evaluation metrics, we should be focusing on the macro average f1-score because we don't want to make false positive prediction and at the same, we want to detect the defaulters.
2. Since NPA (non-performing asset) is a real problem in this industry, it's important we play safe and shouldn't disburse loans to anyone
  - Answer - Below are the most features and their importance while making the prediction. So, these variables can help the managers to identify which are customers who are more likely to pay the loan amount fully,

### Insights:

- 396030 data points, 26 features, 1 label.

- 80% belongs to the class 0: which is loan fully paid.
- 20% belongs to the class 1: which were charged off.
- Loan Amount distribution / media is slightly higher for Charged\_off loanStatus.
- Probability of CHarged\_off status is higher in case of 60 months term.
- Interest Rate mean and media is higher for Charged\_off LoanStatus.
- Probability of Charged\_off LoanStatus is higher for Loan Grades are E, F, G.
- G grade has the highest probability of having defaulter.
- Similar pattern is visible in sub\_grades probability plot.
- Employment Length has overall same probability of Loan\_status as fully paid and defaulter.
- That means Defaulters has no relation with their Employment length.
- For those borrowers who have rental home, has higher probability of defaulters.
- borrowers having their home mortgage and owns have lower probability of defaulter.
- Annual income median is lightly higher for those whos loan status is as fully paid.
- Somehow, verified income borrower's probability of defaulter is higher than those who are not verified by loan tap.
- Most of the borrowers take loans for dept-consolidation and credit card payoffs.
- the probability of defaulters is higher in the small\_business owner borrowers.
- debt-to-income ratio is higher for defaulters.
- number of open credit lines in the borrower's credit file is same as for loan status as fully paid and defaulters.
- Number of derogatory public records increases, the probability of borrowers declared as defaulters also increases
- specially for those who have higher than 12 public\_records.
- Total credit revolving balance is almost same for both borrowers who had fully paid loan and declared defaulter
- but Revolving line utilization rate is higher for defaulter borrowers.
- Application type Direct-Pay has higher probability of defaulter borrowers than individual and joint.
- Number of public record bankruptcies increases, higher the probability of defaulters.



- Most important features/ data for prediction, as per Logistic Regression, Decision tree classifier and Random Forest model are: Employee Title, Loan Grade and Sub-Grade, Interest rate and dept-to-income ratio.